

# Transformers

Prashant Kumar Singh

September 22, 2025

## Abstract

This report documents the implementation and results of Lab 5: Transformers, which focuses on working with pre-trained BERT embeddings, fine-tuning BERT for classification tasks, and extending BERT with a simple autoregressive head. The project demonstrates practical applications of transformer models including text similarity analysis, sentiment classification, and text generation.

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
1.1	Project Overview . . . . .	2
1.2	Technical Objectives . . . . .	2
<b>2</b>	<b>Technical Setup</b>	<b>2</b>
2.1	Environment Configuration . . . . .	2
2.2	Hardware Configuration . . . . .	2
<b>3</b>	<b>Task 1: Text Similarity with BERT Embeddings</b>	<b>2</b>
3.1	Objective . . . . .	2
3.2	Implementation . . . . .	2
3.3	Key Features . . . . .	3
3.4	Results . . . . .	3
<b>4</b>	<b>Task 2: Fine-Tuning RoBERTa for Classification</b>	<b>3</b>
4.1	Objective . . . . .	3
4.2	Data Preparation . . . . .	3
4.3	Model Configuration . . . . .	4
4.4	Training Strategy . . . . .	4
4.5	Results . . . . .	4
<b>5</b>	<b>Task 3: Autoregressive BERT Extension</b>	<b>4</b>
5.1	Objective . . . . .	4
5.2	Implementation . . . . .	4
5.3	Key Features . . . . .	5
5.4	Example Results . . . . .	5
<b>6</b>	<b>Experimental Results</b>	<b>5</b>
6.1	Task 1: Similarity Analysis . . . . .	5
6.2	Task 2: Classification Performance . . . . .	5
6.3	Task 3: Generation Quality . . . . .	5
<b>7</b>	<b>Challenges and Solutions</b>	<b>5</b>
7.1	Technical Challenges . . . . .	5
7.2	Solutions Implemented . . . . .	5
<b>8</b>	<b>Conclusion</b>	<b>6</b>
8.1	Key Achievements . . . . .	6
8.2	Future Work . . . . .	6

# 1 Introduction

## 1.1 Project Overview

This project implements three main tasks using transformer models:

- Text similarity analysis using pre-trained BERT embeddings
- Fine-tuning RoBERTa for text classification
- Autoregressive text generation using BERT with causal masking

## 1.2 Technical Objectives

- Extract contextual embeddings using pre-trained BERT models
- Fine-tune transformer models on downstream classification tasks
- Build simple autoregressive extensions of BERT
- Implement practical NLP applications using Hugging Face transformers

# 2 Technical Setup

## 2.1 Environment Configuration

The project was implemented in Google Colab with GPU acceleration enabled. Key libraries and versions:

```
1 # Core libraries installation
2 !pip install transformers torch scikit-learn datasets matplotlib
3
4 # Specific versions used
5 - transformers: 4.51.3
6 - torch: 2.6.0+cu124
7 - datasets: 3.6.0
8 - scikit-learn: 1.5.2
```

## 2.2 Hardware Configuration

- GPU: NVIDIA Tesla T4 (Google Colab)
- CUDA: Enabled for accelerated training
- Memory: 12GB GPU RAM

# 3 Task 1: Text Similarity with BERT Embeddings

## 3.1 Objective

Implement a function to compute text similarity using cosine similarity between BERT embeddings of different texts.

## 3.2 Implementation

```
1 def text_similarity(
2     texts: Iterable[str],
3     model: BertModel,
4     tokenizer: BertTokenizer,
5     device: torch.device,
6     batch_size: int = 32,
7     max_length: int = 128
8 ) -> np.ndarray:
9     model.eval()
10    model.to(device)
11
```

```

12 all_embeddings = []
13 with torch.no_grad():
14     for i in range(0, len(texts), batch_size):
15         batch_texts = texts[i : i + batch_size]
16         encoded = tokenizer(
17             batch_texts,
18             padding='max_length',
19             truncation=True,
20             max_length=max_length,
21             return_tensors='pt'
22         ).to(device)
23
24         outputs = model(**encoded)
25         cls_emb = outputs.pooler_output
26         all_embeddings.append(cls_emb.cpu().numpy())
27
28 embeddings = np.vstack(all_embeddings)
29 similarity = cosine_similarity(embeddings)
30 return similarity.astype(np.float32)

```

### 3.3 Key Features

- Batch processing for memory efficiency
- Uses BERT's pooler output for sentence embeddings
- Cosine similarity matrix computation
- GPU acceleration support

### 3.4 Results

The function successfully computed similarity matrices for sentences in `data/task1/sentences.csv`. The output was saved as `similarity.csv` for submission.

## 4 Task 2: Fine-Tuning RoBERTa for Classification

### 4.1 Objective

Fine-tune a RoBERTa model to classify text data, achieving F1-score  $\geq 0.77$ .

### 4.2 Data Preparation

```

1 class TextDataset(Dataset):
2     def __init__(self, file_paths, labels=None, tokenizer=None, max_length=512):
3         self.files = list(file_paths)
4         self.labels = labels
5         self.tokenizer = tokenizer
6         self.max_length = max_length
7
8     def __getitem__(self, idx):
9         path = self.files[idx]
10        with open(path, 'r', encoding='utf-8') as f:
11            text = f.read()
12            enc = self.tokenizer(text, truncation=True, padding='max_length',
13                                max_length=self.max_length, return_tensors='pt')
14            item = {k: v.squeeze(0) for k, v in enc.items()}
15            if self.labels is not None:
16                item['labels'] = torch.tensor(self.labels[idx], dtype=torch.long)
17        return item

```

## 4.3 Model Configuration

```
1 # Model initialization
2 model = RobertaForSequenceClassification.from_pretrained(
3     'roberta-base',
4     num_labels=2
5 ).to(device)
6
7 # Optimizer and scheduler
8 optimizer = AdamW(model.parameters(), lr=1e-5)
9 scheduler = get_linear_schedule_with_warmup(
10     optimizer,
11     num_warmup_steps=int(0.1 * total_steps),
12     num_training_steps=total_steps
13 )
```

## 4.4 Training Strategy

- Learning rate: 1e-5 (low for fine-tuning)
- Batch size: 8
- Early stopping with patience=2
- Stratified train-validation split (90%-10%)
- Linear learning rate scheduling with warmup

## 4.5 Results

The model achieved the target F1-score  $\geq 0.77$ . Predictions were saved to `submission.csv`.

# 5 Task 3: Autoregressive BERT Extension

## 5.1 Objective

Implement a simple autoregressive text generation system using BERT with causal masking.

## 5.2 Implementation

```
1 def complete_text(prompt: str, max_tokens: Optional[int] = None,
2                   model=model, tokenizer=tokenizer, device=device):
3     if max_tokens is None:
4         max_tokens = tokenizer.model_max_length - len(tokenizer(prompt).input_ids)
5
6     prompt += ' '.join(['[MASK]'] * max_tokens)
7     inputs = tokenizer(prompt, return_tensors="pt").to(device)
8
9     with torch.no_grad():
10         logits = model(**inputs).logits
11
12     predicted_token_id = logits[0, -max_tokens-1].argmax(axis=-1).cpu().tolist()
13     text = tokenizer.decode(inputs.input_ids[0, 1:-max_tokens-1].cpu().tolist() +
14                             [predicted_token_id])
15
16     # Stopping conditions
17     if (max_tokens == 1 or
18         predicted_token_id == tokenizer.vocab['.'] or
19         predicted_token_id == tokenizer.sep_token_id):
20         return text
21
22     return complete_text(text, max_tokens=max_tokens-1,
23                           model=model, tokenizer=tokenizer, device=device)
```

### 5.3 Key Features

- Recursive text generation
- Automatic stopping on punctuation
- Mask token utilization for next-word prediction
- GPU-accelerated inference

### 5.4 Example Results

Input: "The capital of France is "  
Output: "The capital of France is Paris."

Input: "one plus one is equal to "  
Output: "one plus one is equal to two."

## 6 Experimental Results

### 6.1 Task 1: Similarity Analysis

The text similarity function successfully generated cosine similarity matrices demonstrating semantic relationships between sentences.

### 6.2 Task 2: Classification Performance

Metric	Training	Validation
F1-Score	0.85+	0.77
Accuracy	0.87+	0.79
Loss	0.24	0.28

Table 1: Classification performance metrics

### 6.3 Task 3: Generation Quality

The autoregressive BERT extension demonstrated coherent text generation capabilities, though limited by BERT’s non-autoregressive architecture.

## 7 Challenges and Solutions

### 7.1 Technical Challenges

1. **Memory Management:** Large transformer models require careful batch sizing
2. **Training Stability:** Low learning rates and scheduling prevent overfitting
3. **Text Generation:** BERT’s architecture limitations for autoregressive tasks

### 7.2 Solutions Implemented

- Gradient accumulation and mixed precision training
- Early stopping and learning rate scheduling
- Custom stopping conditions for text generation

## 8 Conclusion

This project successfully demonstrated practical applications of transformer models across three distinct tasks:

### 8.1 Key Achievements

- Implemented efficient text similarity analysis using BERT embeddings
- Achieved target performance in text classification through careful fine-tuning
- Developed working autoregressive text generation with BERT
- Established reproducible training pipelines with proper evaluation

### 8.2 Future Work

- Experiment with larger transformer architectures (BERT-large, GPT-2)
- Implement more sophisticated generation techniques (beam search, sampling)
- Explore multi-lingual transformer applications
- Optimize for production deployment scenarios